



University
of Glasgow



Fjelltopp
Technology with impact.

session 1:
Basic Predictive Models

M. Kundegorski

3rd March 2020

Centre for Ecological Sciences
Indian Institute of Science
Bengaluru, India

Basic Models

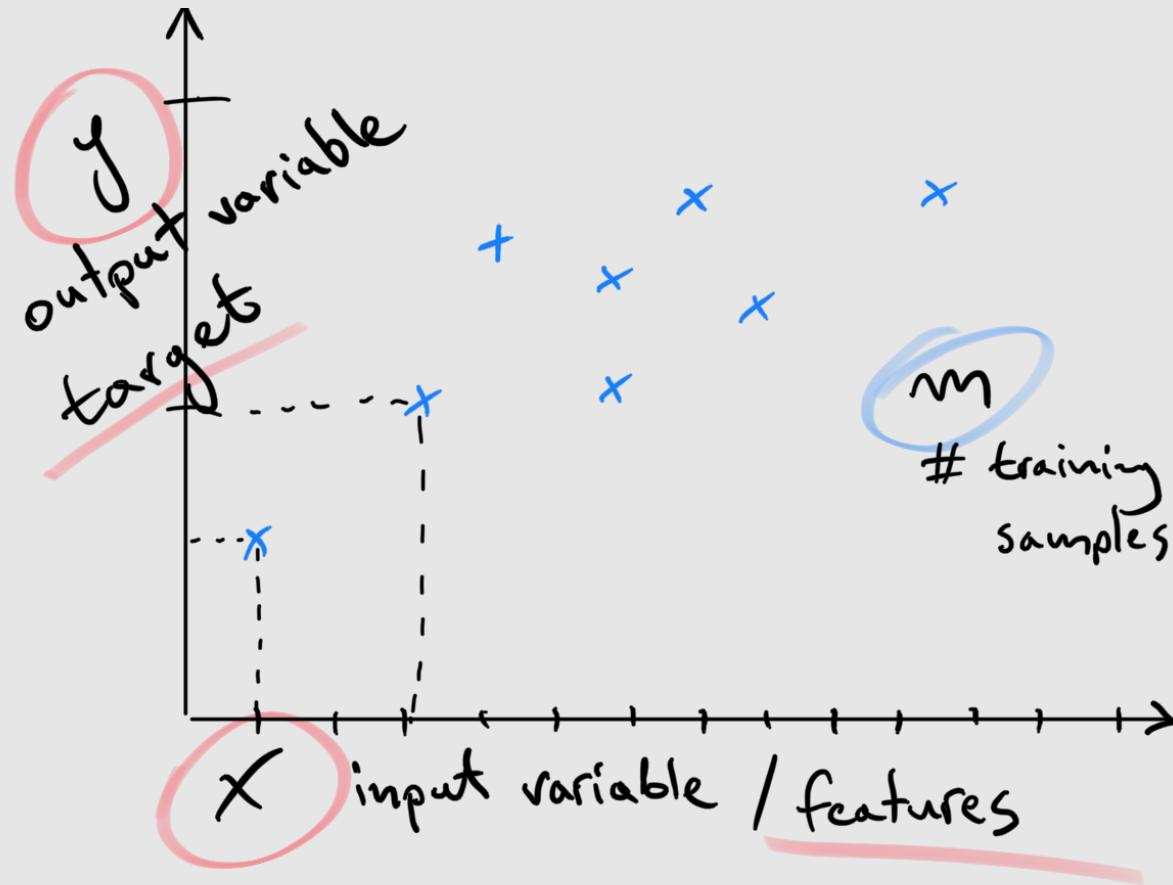
Statistics from a different point of view

In this lecture

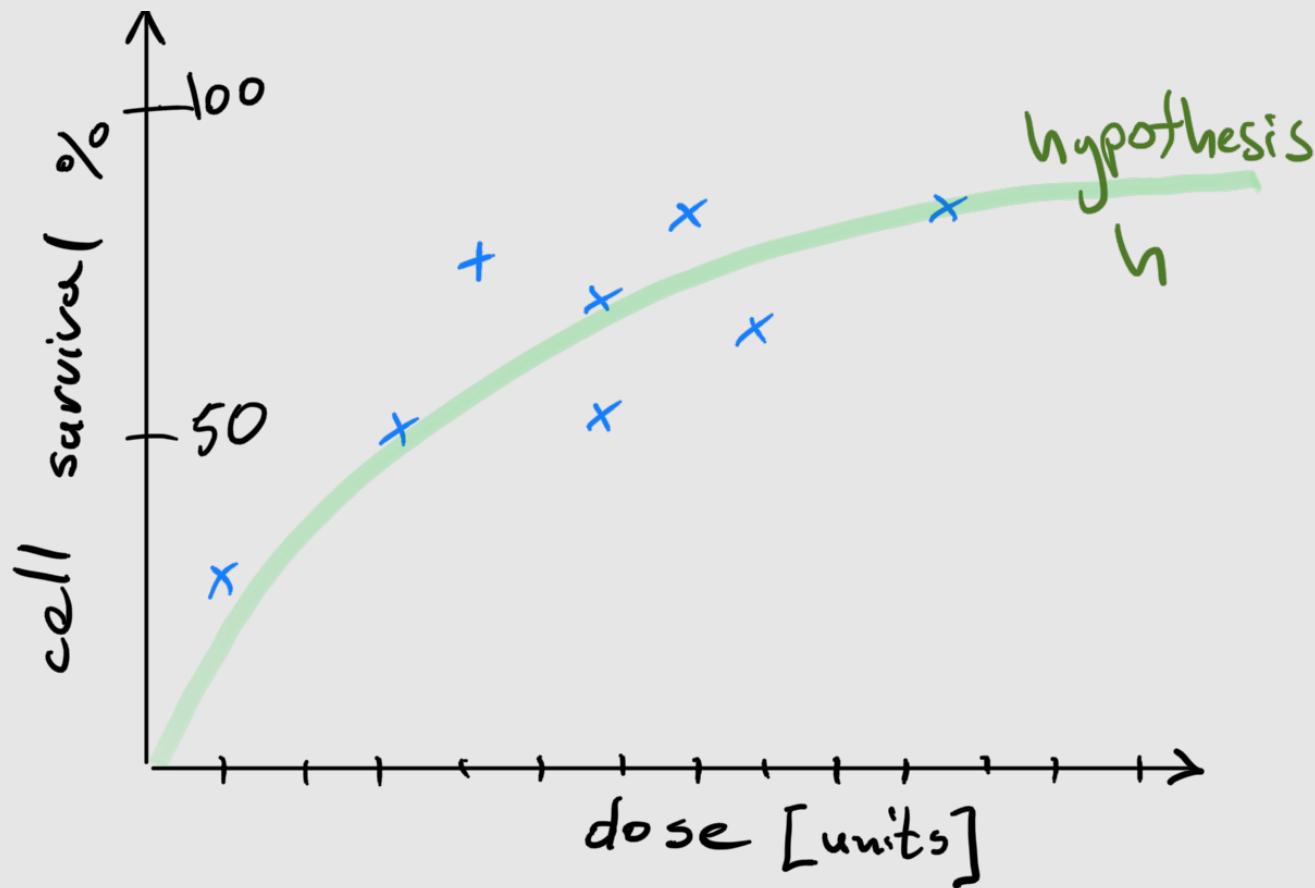
- Linear regression
- Multivariate regression
- Generalised linear regression
- Model fitting

Based on excellent online course by Andrew Ng:
<https://www.coursera.org/learn/machine-learning/>

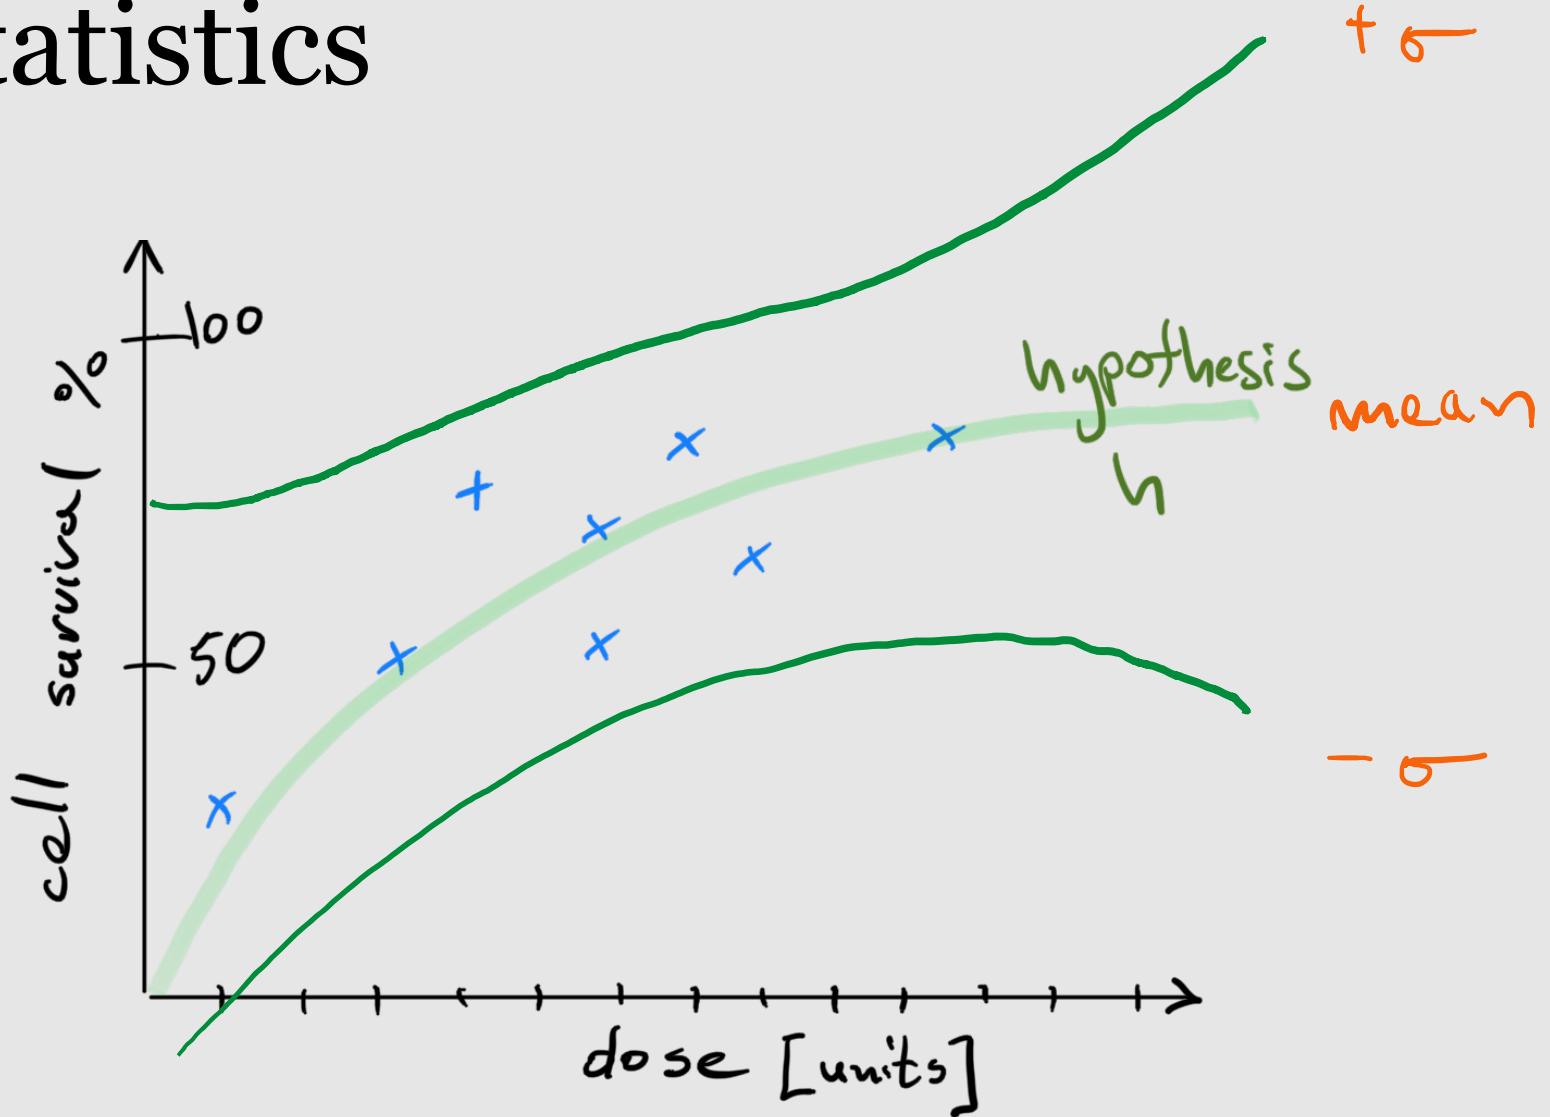
A problem



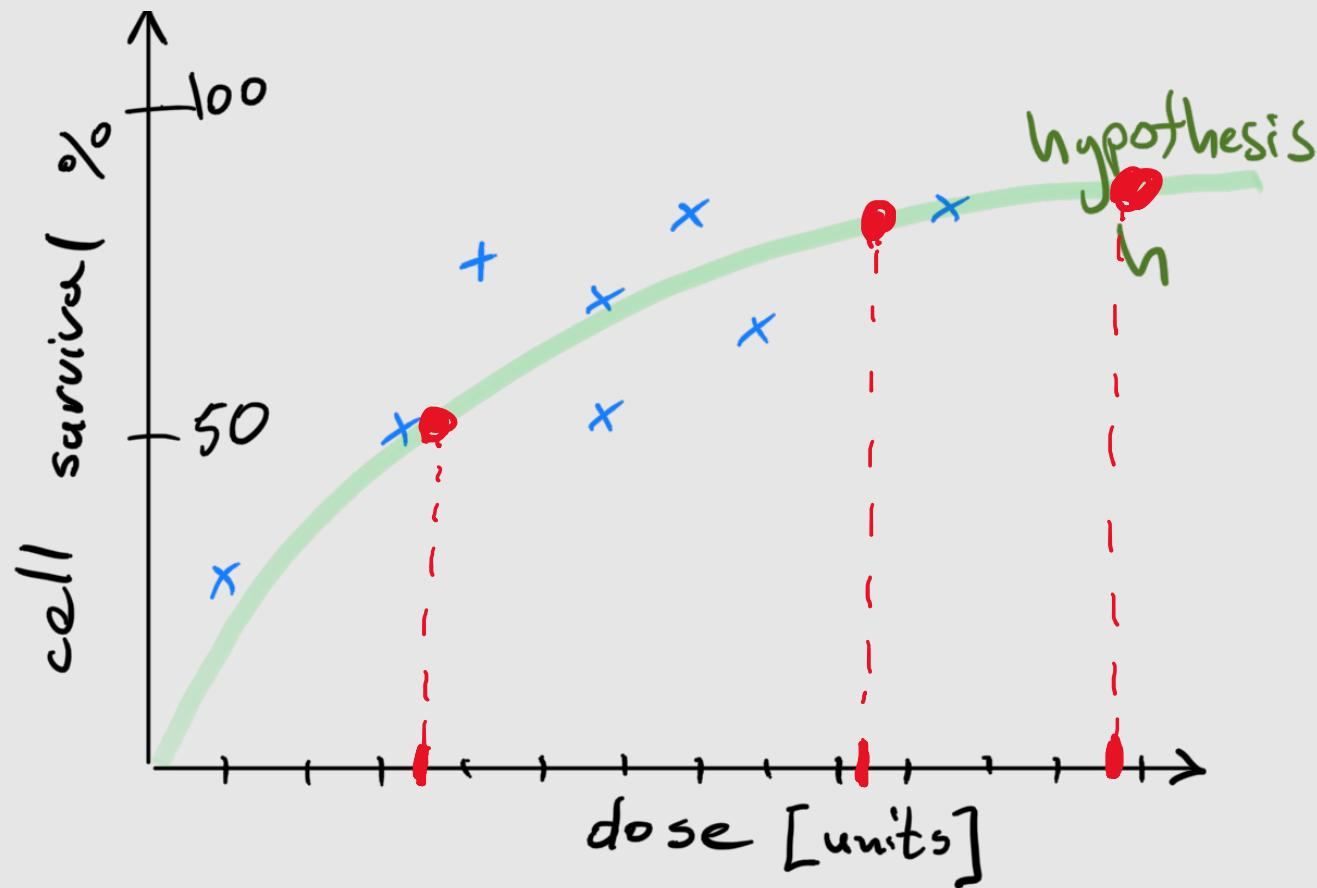
Model Representation



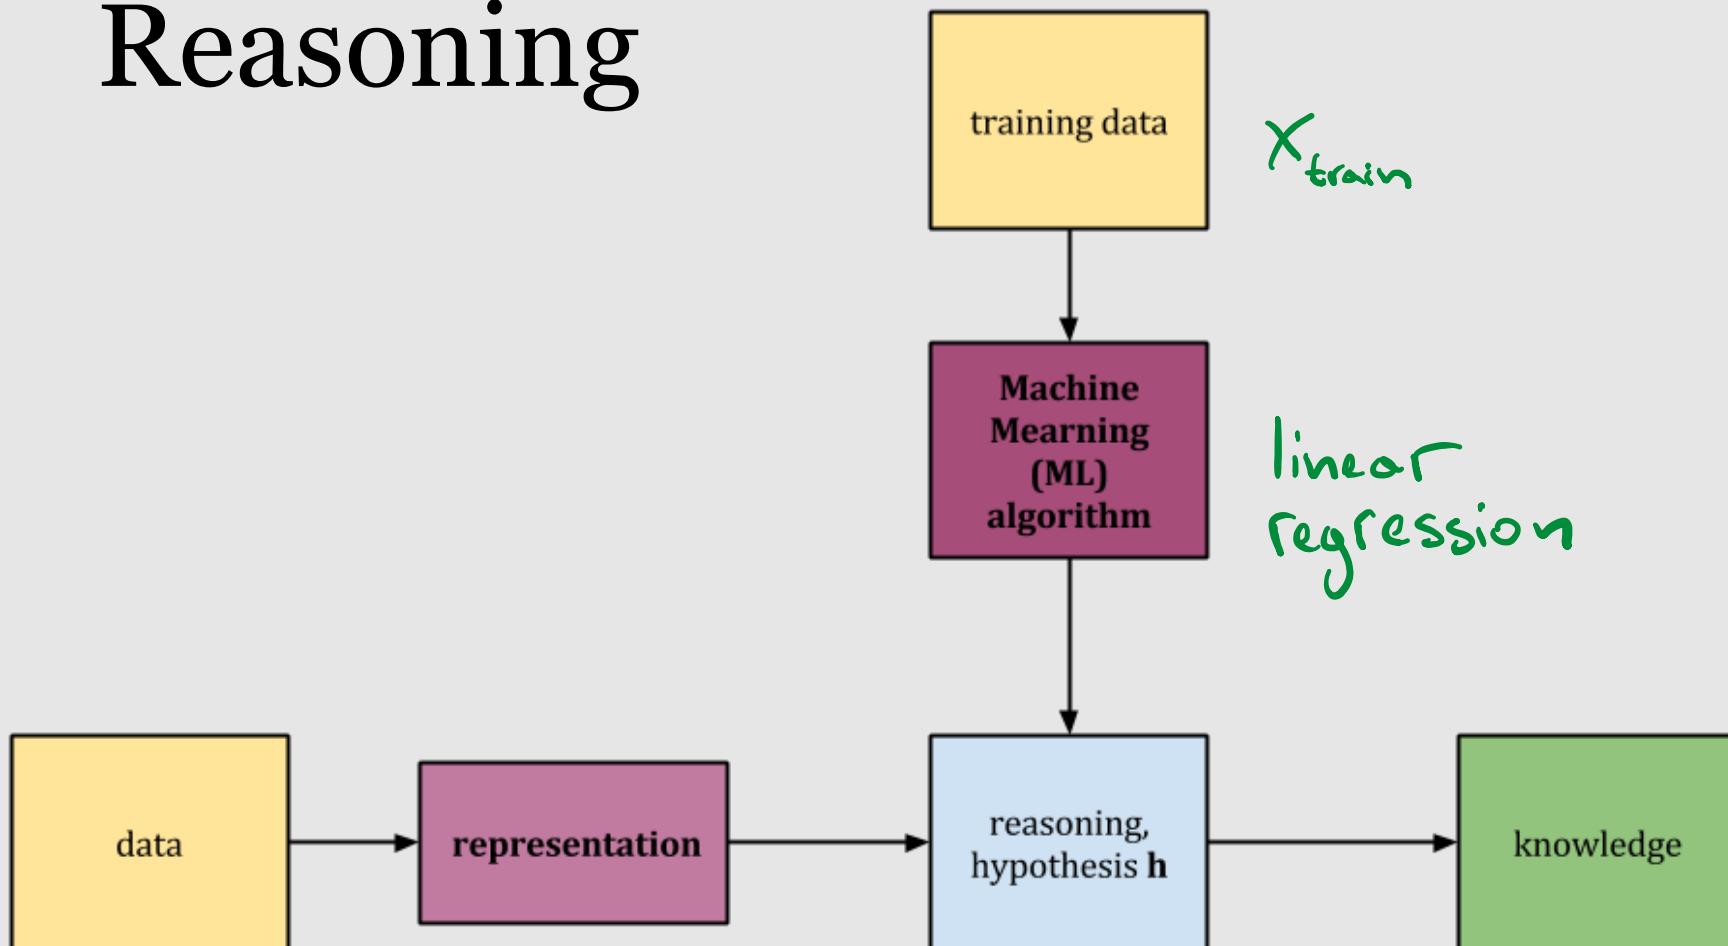
Statistics



Machine learning

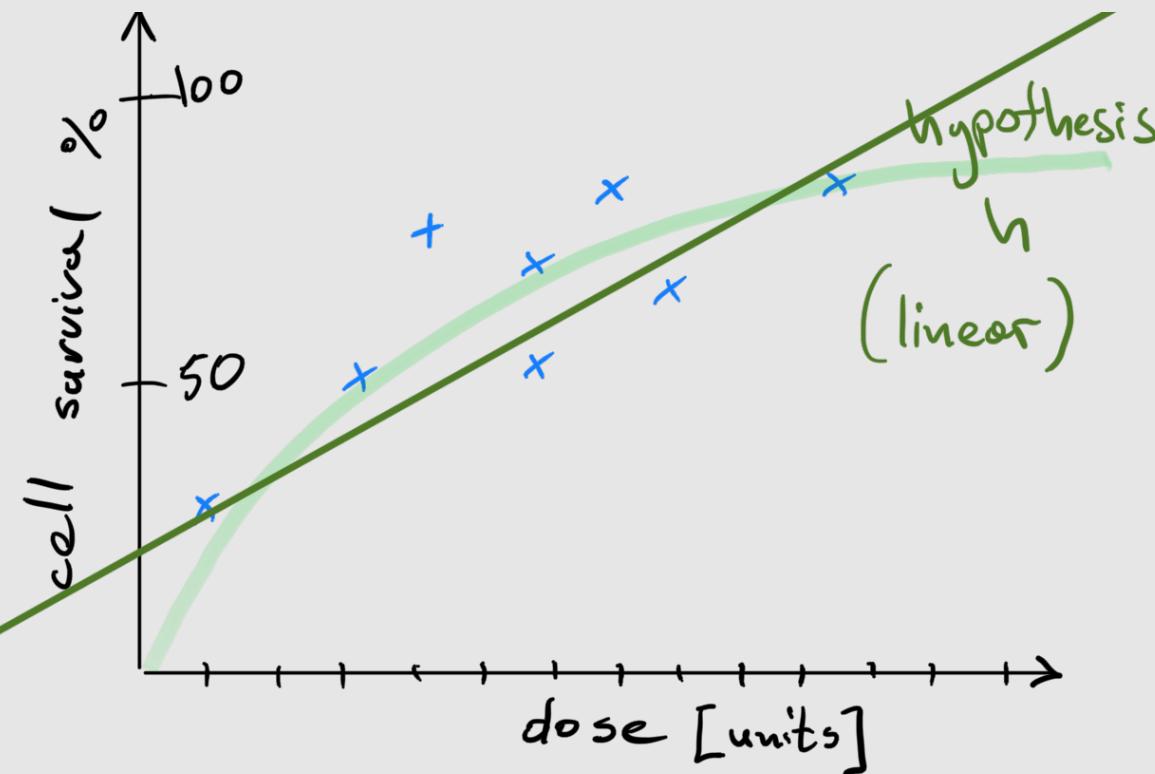


Reasoning



$$X_? \rightarrow h_{\theta}(x) = \theta_0 + \theta_1 x \rightarrow y_{predicted}$$

Linear regression

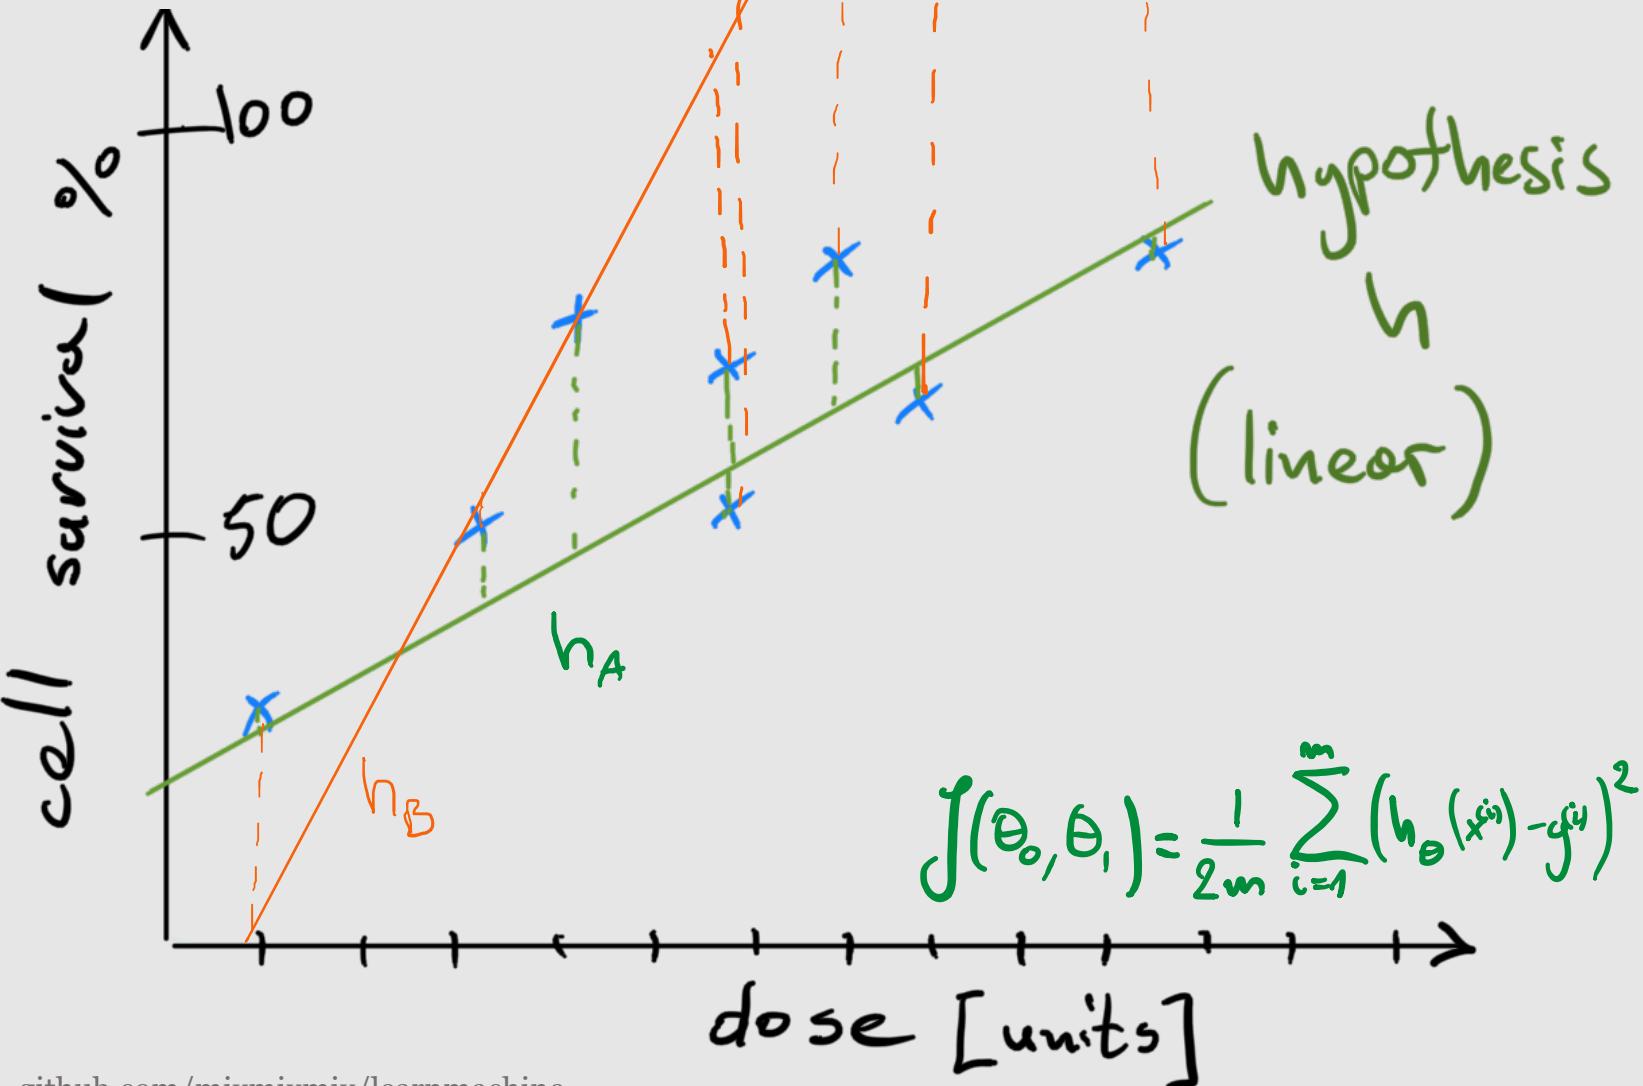


$$\frac{1}{2m} \sum_{i=0}^m (h_0(x_i) - y_i)^2$$

Squared error cost function.

We have m training examples
 (x_i, y_i)

Fitting error



Cost Function

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

Parameters: θ_0, θ_1

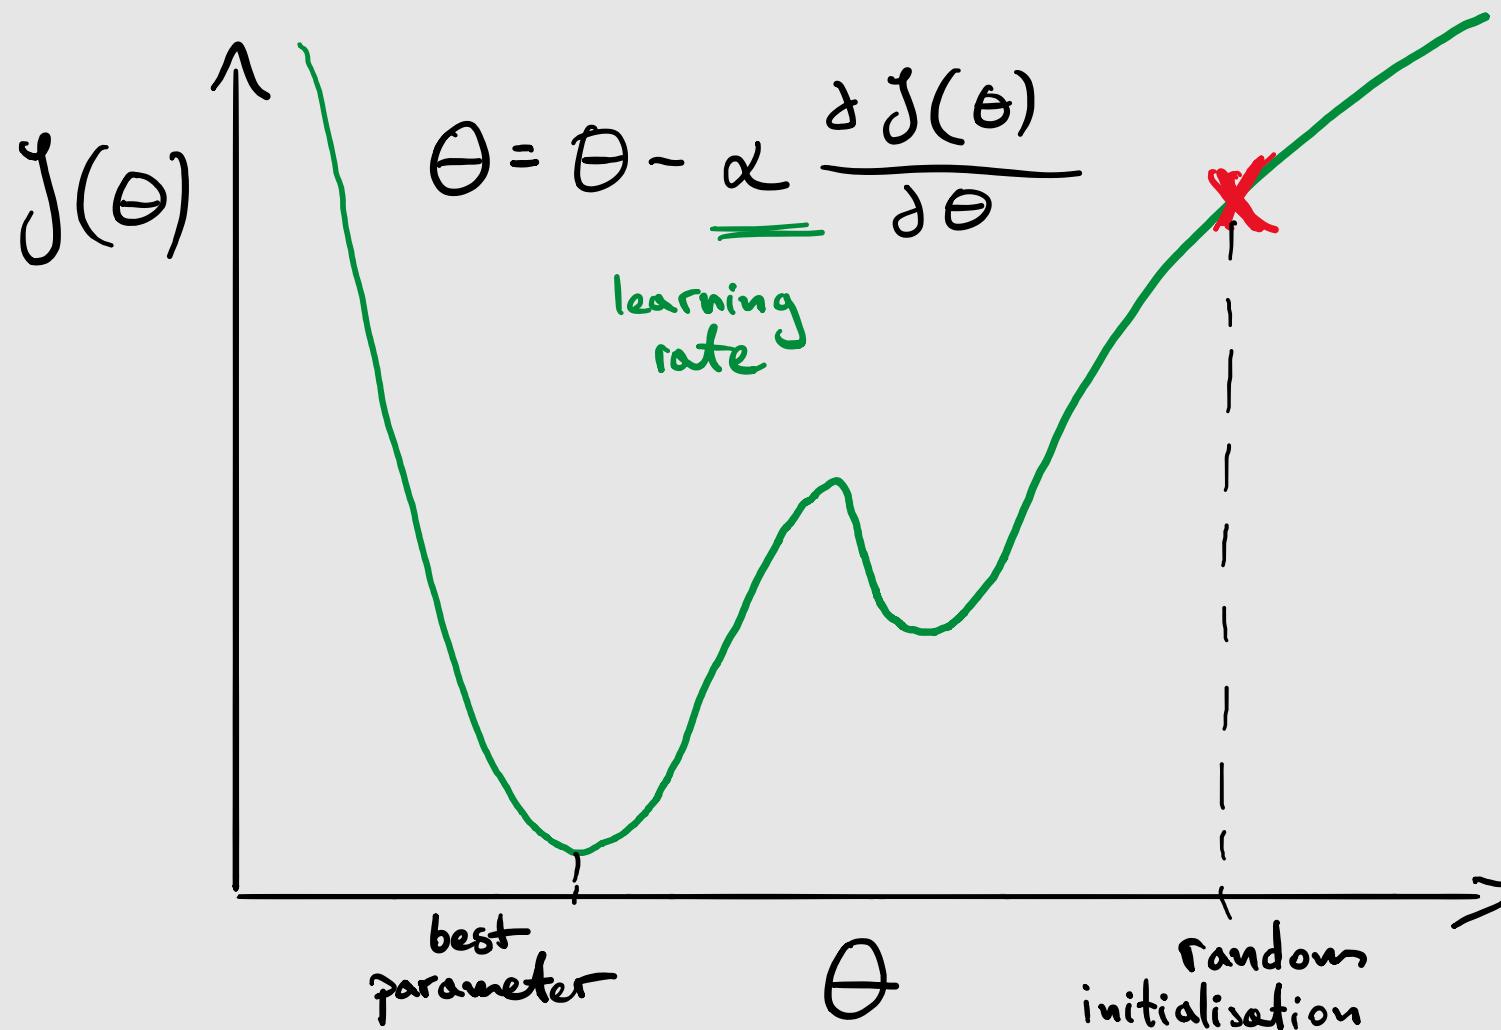
Cost function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Goal: $\text{argmin } J(\theta_0, \theta_1)$

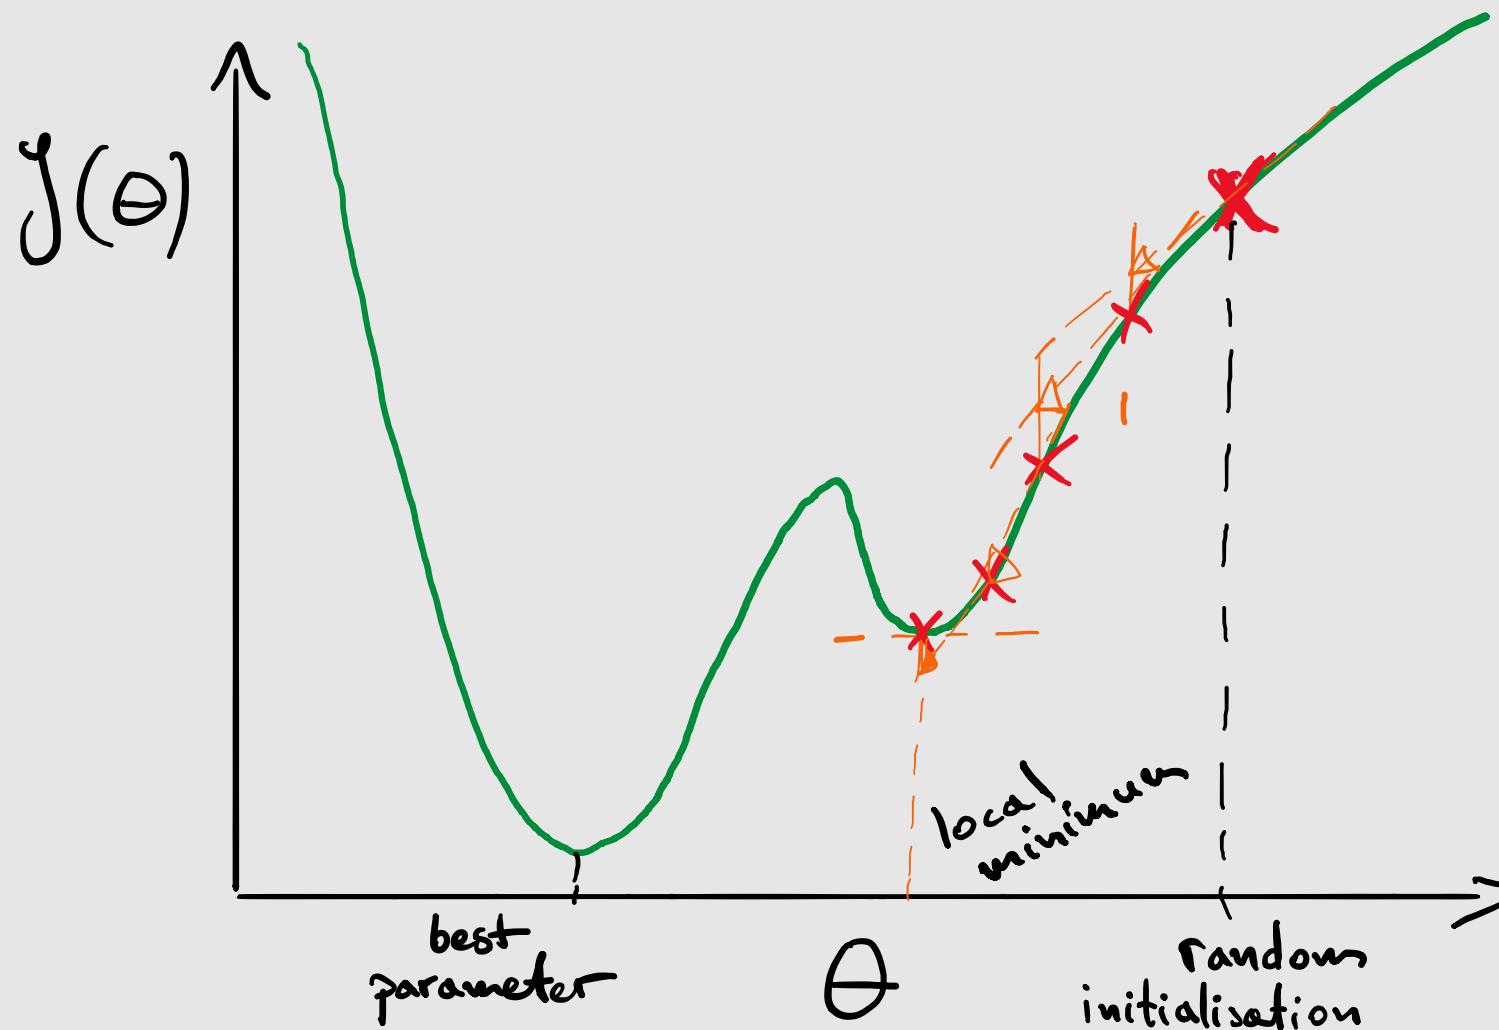


aka *Loss function, Error function*

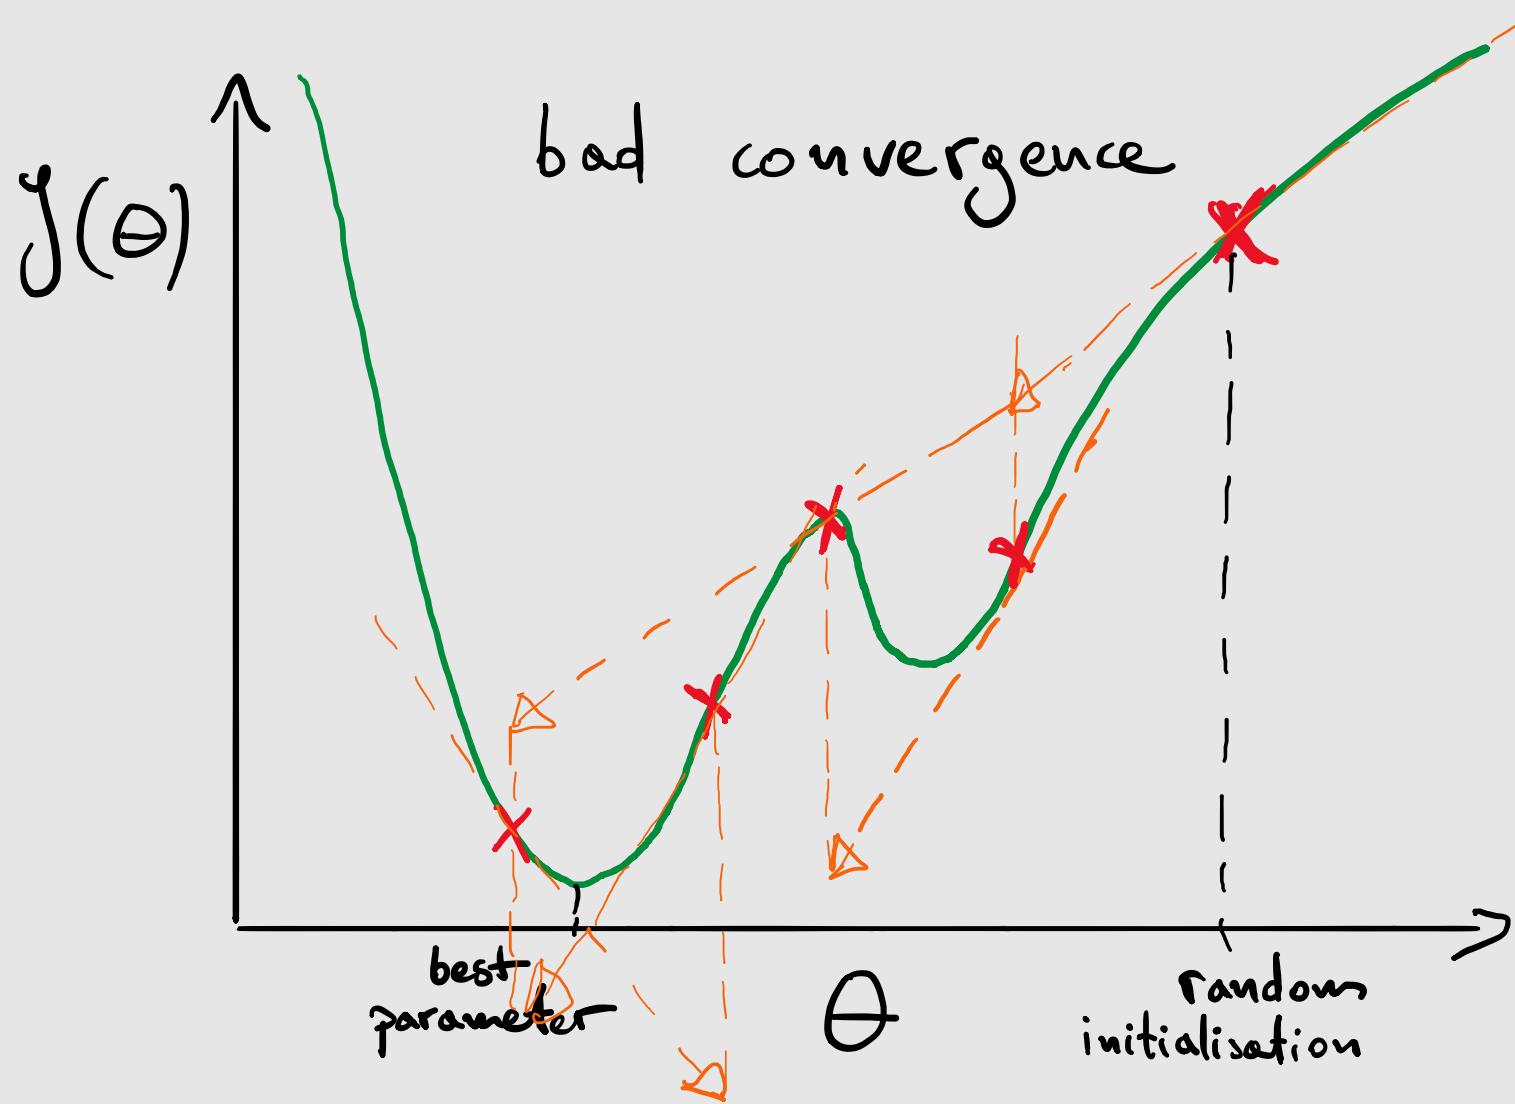
Gradient descent



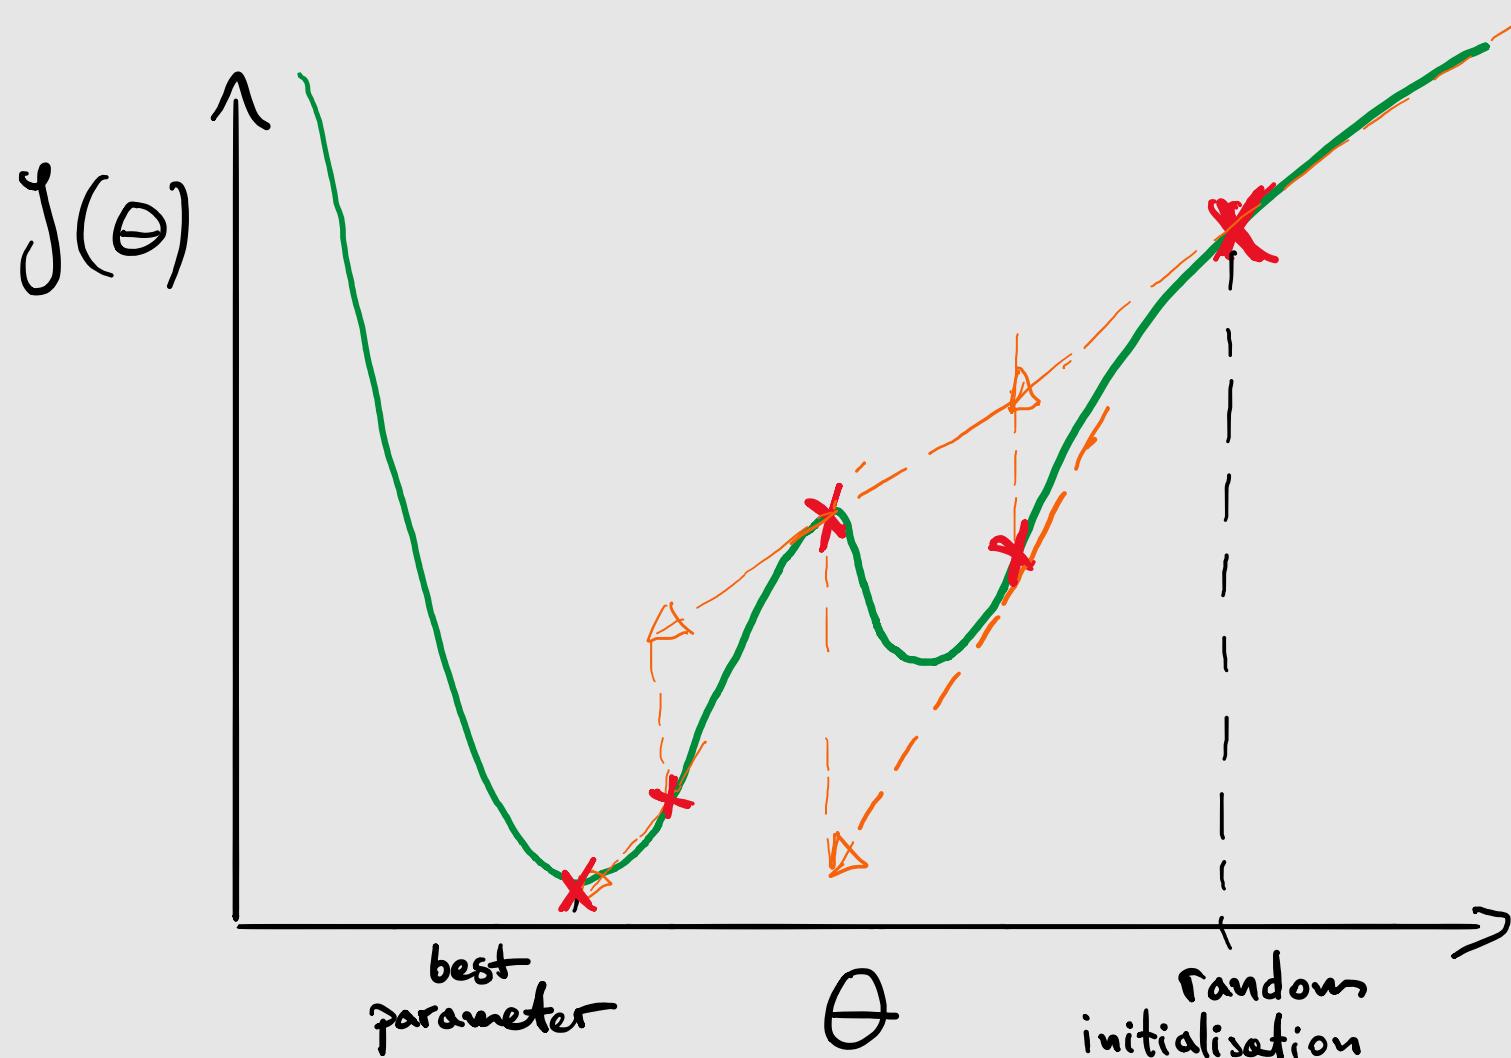
Gradient descent (small α)



Gradient descent (big α)



Gradient descent (adaptive α)



Gradient descent

- all parameters simultaneous
- batch: uses a "batch" of training examples

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x_i) - y_i)^2$$

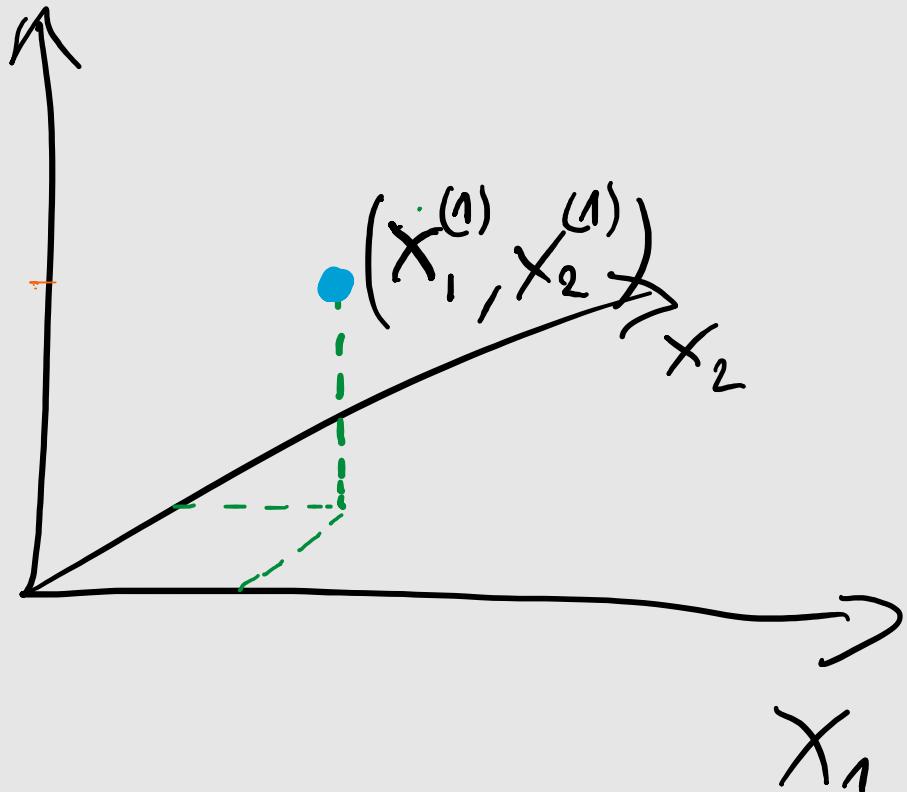
$$\frac{\partial}{\partial \theta_j} J(\theta) = (h_\theta(x) - y) x_j$$

Multiple features

$$h_{\theta}(x) = \theta_0 \overset{x_0=1}{\nearrow} x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

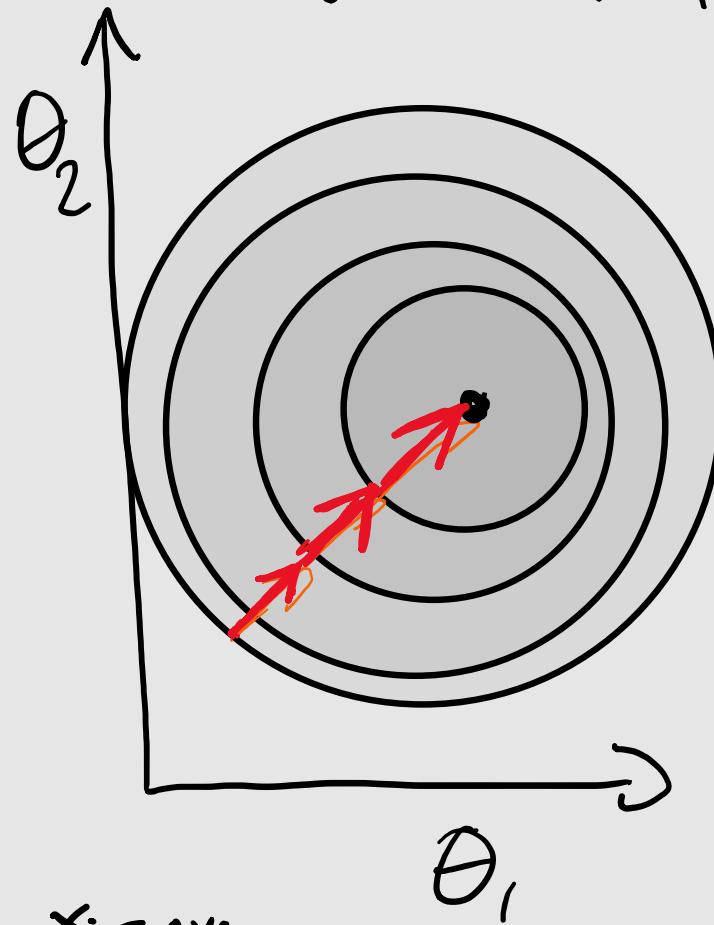
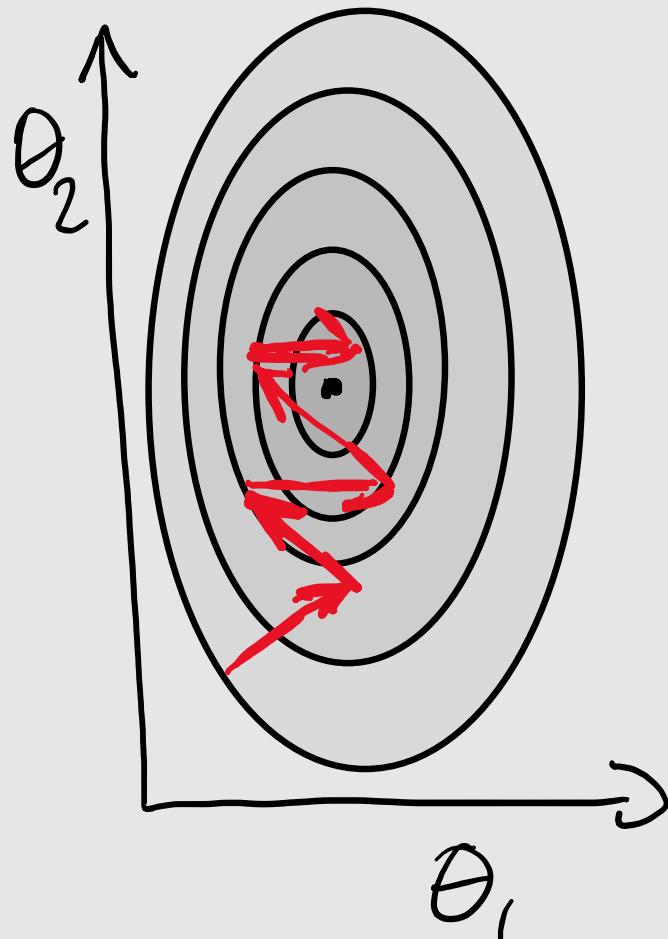
$$h_{\theta}(x) = \theta^T x$$

$$h_{\theta}(x) = [\theta_0, \theta_1, \dots, \theta_n] \begin{bmatrix} 1 \\ x_1^{(1)} \\ x_2^{(1)} \\ \vdots \\ x_n^{(1)} \end{bmatrix}$$



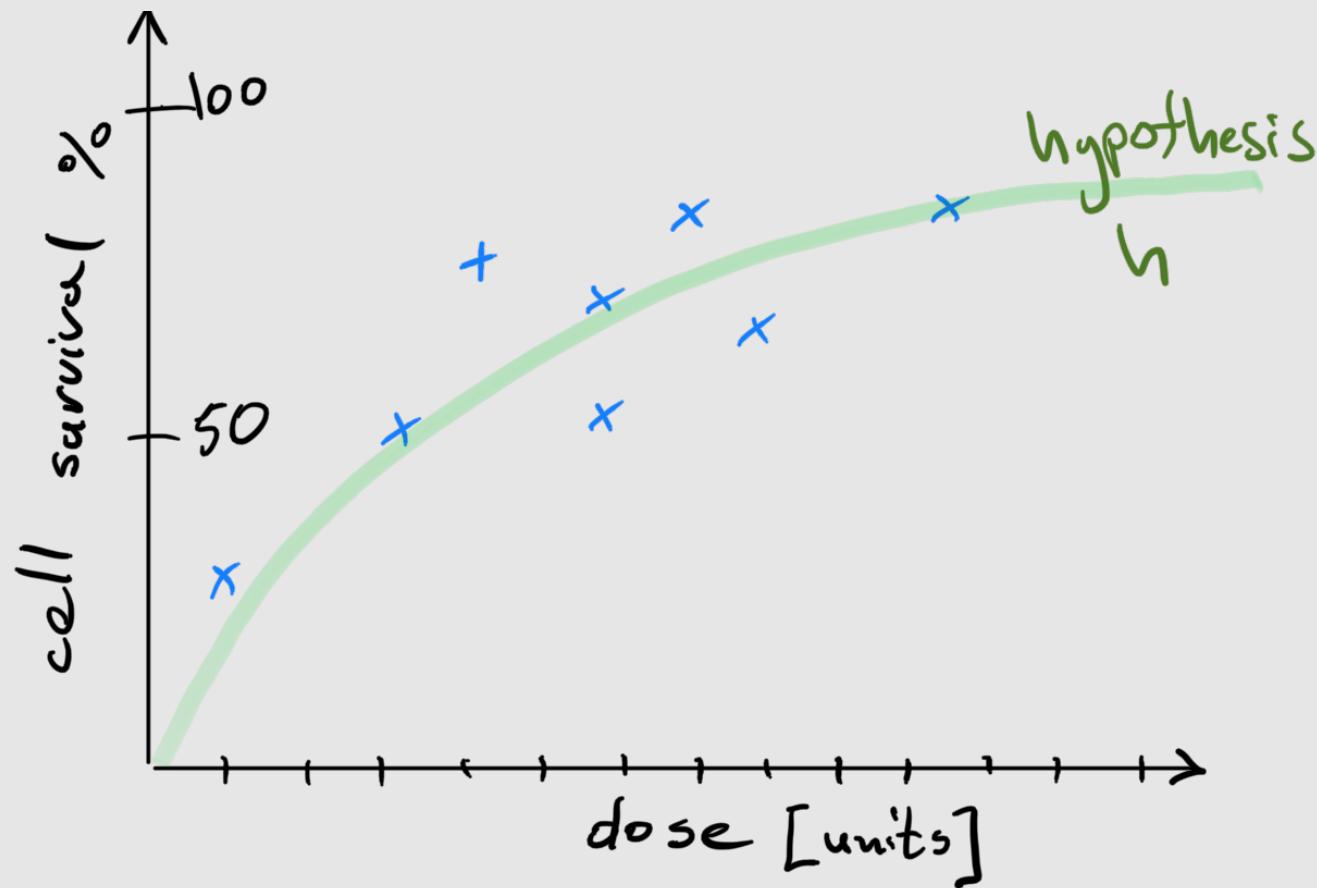
Multiple features

α - learning rate doesn't like un-scaled features



$$x_{\text{scaled}} = \frac{x_i - \text{avg}}{\text{max} - \text{min}}$$

Non-linear regression

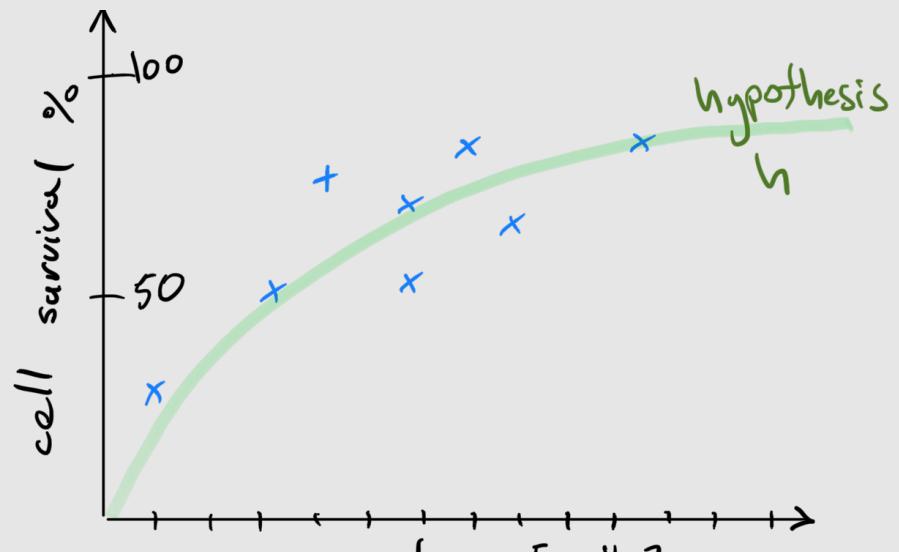


Multiple non-linear features

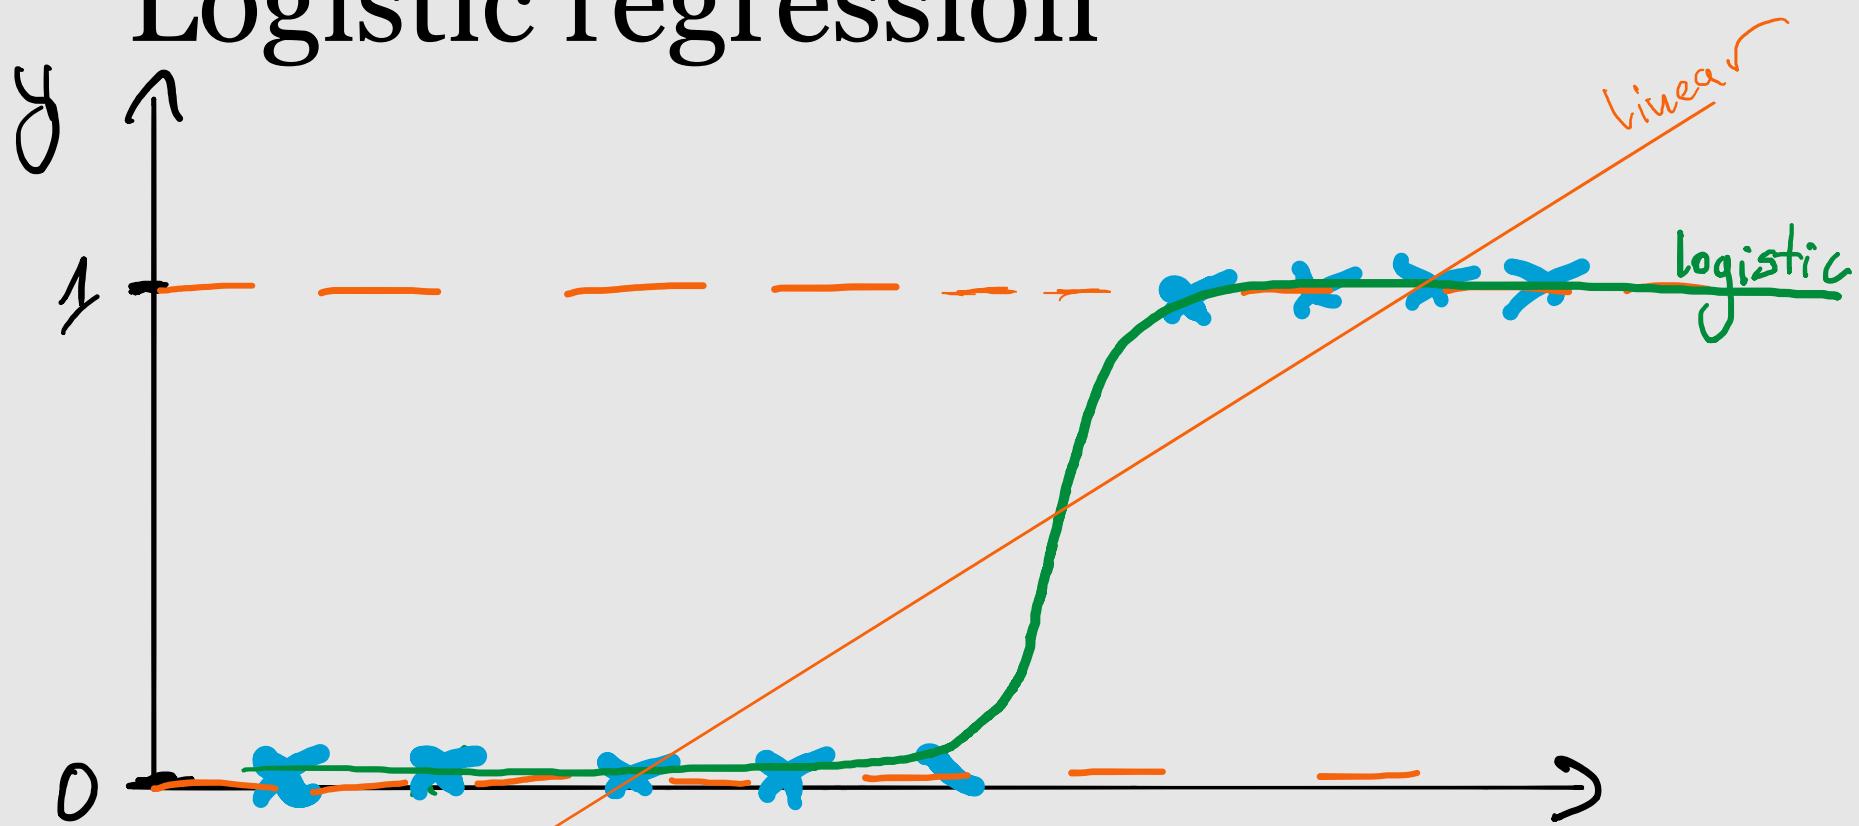
$$h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \dots + \theta_n x^n \quad \left. \right\} \text{polynomial}$$

$$h_{\theta}(x) = \theta_0 + \theta_1 \sqrt{x} + \theta_2 \log(x) + \dots \quad \left. \right\} \text{non-linear (generalised)}$$

$J(\theta)$ → stays the same.



Logistic regression



$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

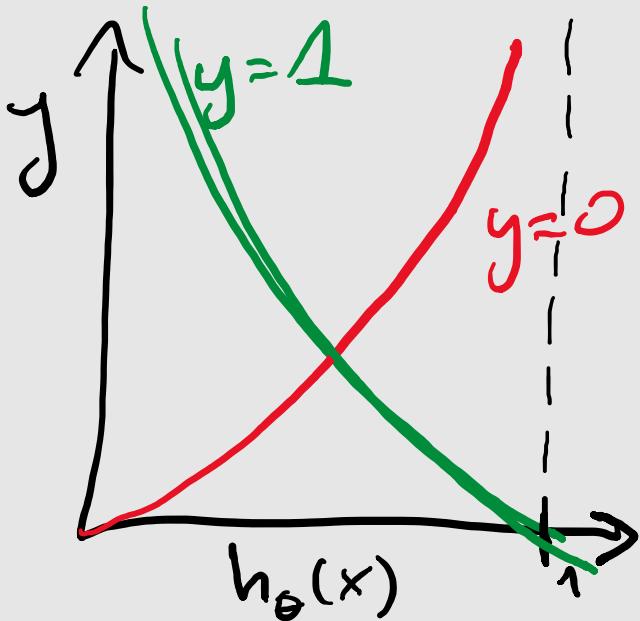
$$J(\theta) = ?$$

Logistic regression

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

normal
sum of squares
makes J
non-convex



$$\text{Cost}(h_{\theta}(x), y) =$$

$$\begin{cases} -\log(h_{\theta}(x)) & : y=1 \\ -\log(1-h_{\theta}(x)) & : y=0 \end{cases}$$

Scikit-learn library

scikit-learn

Machine Learning in Python

Getting Started

What's New in 0.22

GitHub

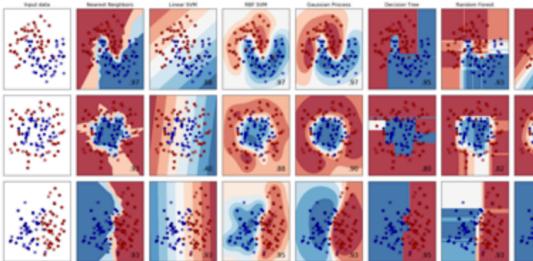
- Simple and efficient tools for predictive data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying which category an object belongs to.

Applications: Spam detection, image recognition.

Algorithms: SVM, nearest neighbors, random forest, and more...



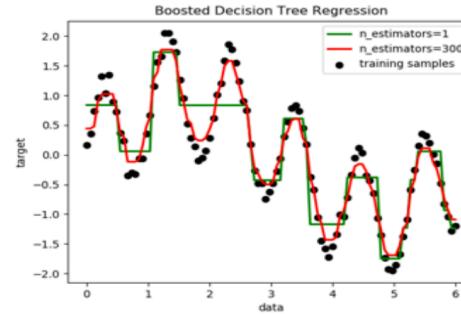
Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, nearest neighbors, random forest, and more...



Examples

Clustering

Automatic grouping of similar objects into sets.

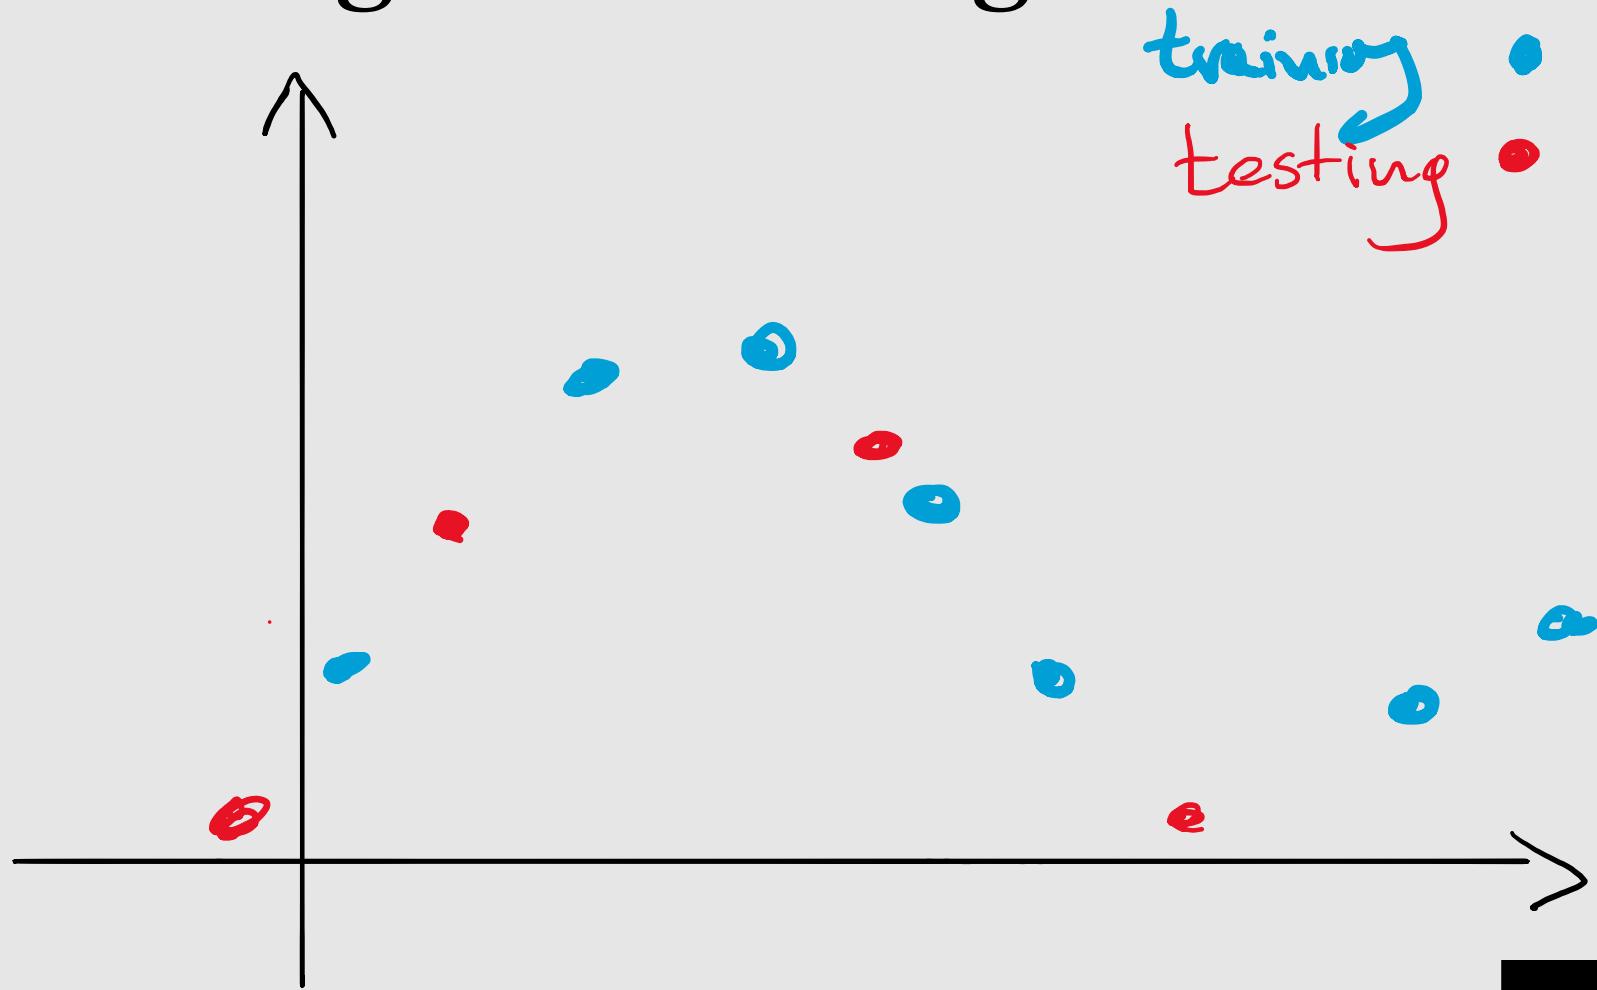
Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, and more...

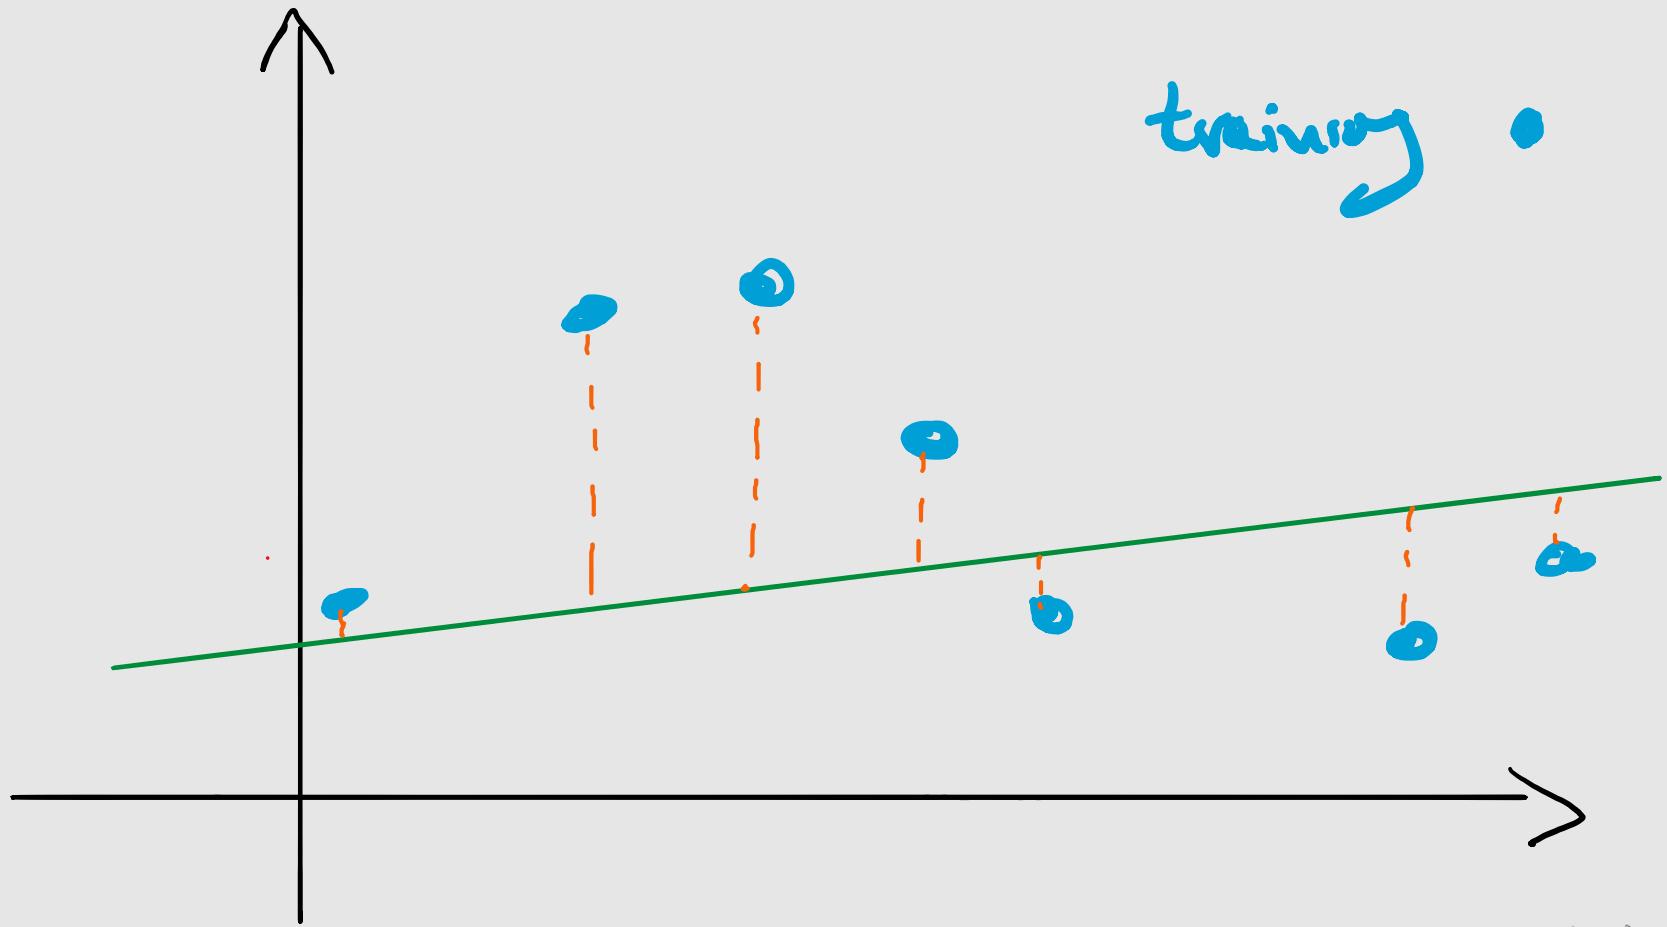


Examples

Training and Testing: data



fit()



predict()

